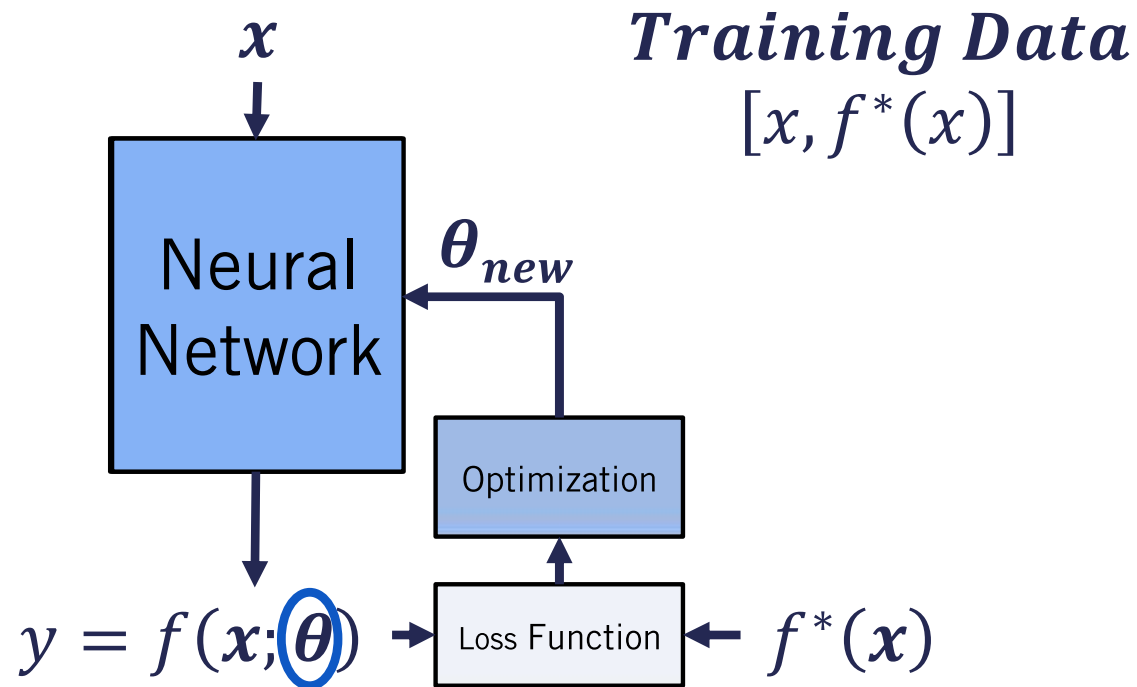


# Learning Objectives

- Learn how to train a neural network using the iterative optimization algorithm: gradient descent
- Learn how to initialize parameters at the start of the optimization process

# Artificial Neural Networks



# Neural Network Loss Functions

- **Thousands** of training example pairs  $[\mathbf{x}, f^*(\mathbf{x})]$
- The **Loss function** computed over all **N** training examples is termed the **Training Loss** and can be written as:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N L[f(x_i, \theta), f^*(x_i)]$$

- The gradient of the training loss with respect to the parameters  $\theta$  can be written as:

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \left[ \frac{1}{N} \sum_{i=1}^N L[f(x_i, \theta), f^*(x_i)] \right] = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)]$$

# Batch Gradient Descent

- Batch Gradient Descent is an **iterative first order** optimization procedure

- **Batch Gradient Descent Algorithm:**

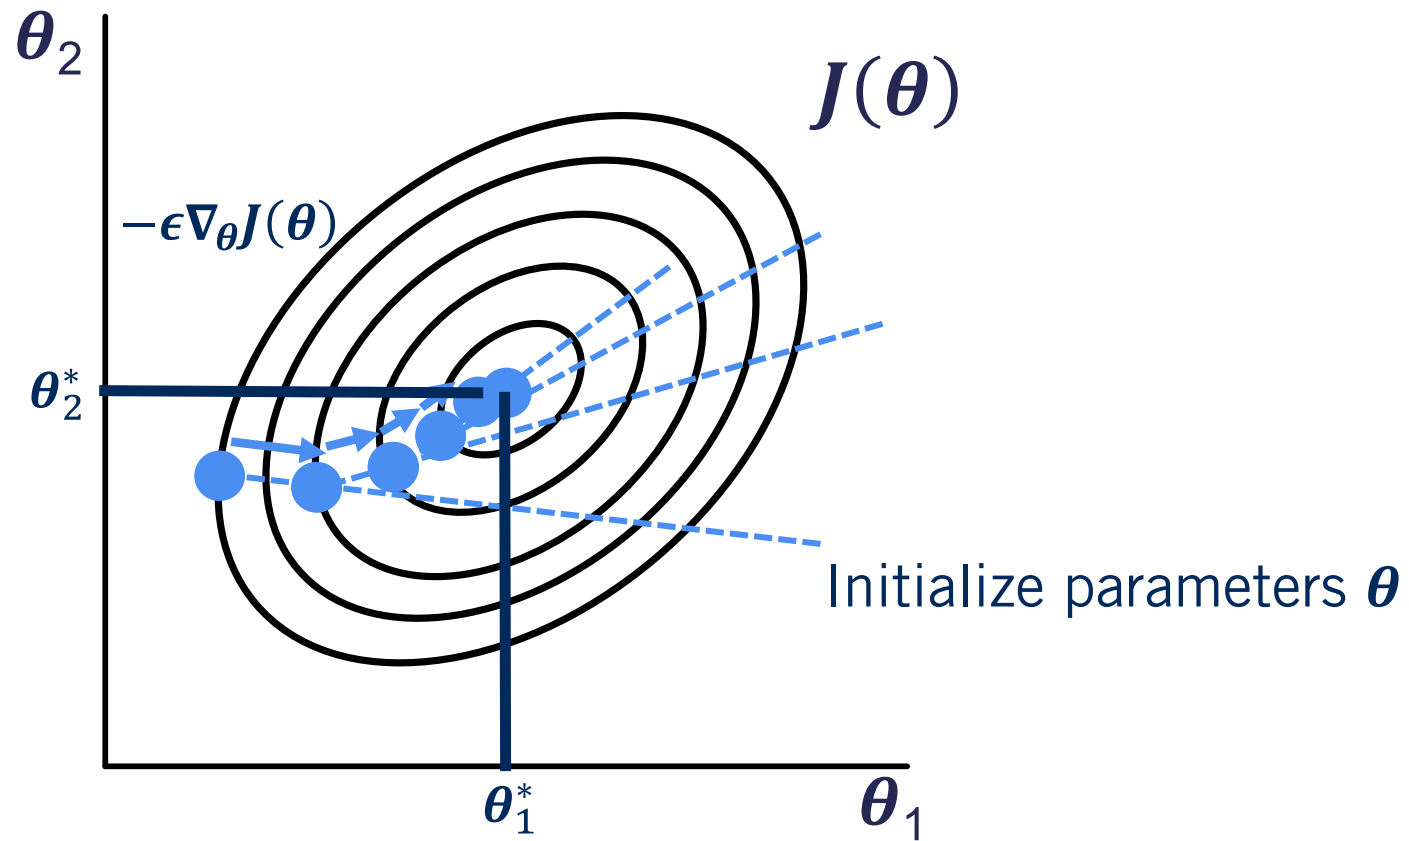
- Initialize parameters  $\theta$
- While **Stopping Condition** is **Not Met**:
  - Compute gradient of loss function over **all** training examples:

$$\nabla_{\theta} J(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} L[f(x_i, \theta), f^*(x_i)]$$

- Update parameters according to:

$$\theta \leftarrow \theta - \epsilon \nabla_{\theta} J(\theta)$$

# Batch Gradient Descent



# Parameter Initialization and Stopping Conditions

- **Parameter Initialization:**

- **Weights:** initialized by randomly sampling from a standard normal distribution
- **Biases:** initialized to 0
- Other **heuristics** exist

- **Stopping Conditions:**

- **Number of iterations:** How many training iterations the neural network has performed
- **Change In  $\theta$  value:** Stop if  $\theta_{new} - \theta_{old} < \text{Threshold}$
- **Change In  $J(\theta)$  value:** Stop if  $J(\theta_{new}) - J(\theta_{old}) < \text{Threshold}$

# Batch Gradient Descent

- **Backpropagation** used to compute  $\nabla_{\theta}J(\theta)$  is very expensive to compute over the whole training dataset.

$$\nabla_{\theta}J(\theta) = \frac{1}{N} \sum_i^N \nabla_{\theta}L [f_i(x, \theta), f_i^*(x)] \text{ is a mean!}$$

- Standard error of the mean estimated from  $N$  samples is  $\frac{\sigma}{\sqrt{N}}$ , where  $\sigma$  is the standard deviation of the value of the samples
- Using all samples to estimate the gradient results in **less than linear** return in accuracy of this estimate
- Use a **small subsample (Minibatch)** of the training data to estimate the gradient!

# Stochastic (minibatch) Gradient Descent

- **Stochastic (minibatch) Gradient Descent Algorithm:**

- Initialize parameters  $\theta$
- While **Stopping Condition** is **False**:
  1. Sample a preset number  $N'$  of examples (minibatch) from the training data
  2. Compute gradient of Loss Function Over **all** training examples:

$$\nabla_{\theta} J(\theta) = \frac{1}{N'} \sum_i^{N'} \nabla_{\theta} L [f_i(x, \theta), f_i^*(x)]$$

3. Update parameters according to:

$$\theta \leftarrow \theta - \epsilon \nabla_{\theta} J(\theta)$$



# What Minibatch Size To Use ?

- GPUs work better with **powers of 2** batch sizes
- **Large batch sizes > 256:**
  - Hardware underutilized with very small batch sizes.
  - More accurate estimate of the gradient, but with less than linear returns
- **Small batch size < 64**
  - Small batches can offer a **regularizing effect**. The best generalization error is often achieved with batch size of 1.
  - Small batch sizes allow for faster convergence, as the algorithm can compute the parameter updates rapidly
- Always make sure dataset is **shuffled** before sampling minibatch

# SGD Variations

- Many variations of SGD exist
  - Momentum SGD, Nesterov Momentum SGD
  - Ada-Grad, RMS-Prop
  - ADAM (Adaptive Moment Estimation)
- Which one to use ?
  - ADAM: Implemented in most deep neural network libraries, fairly robust to the choice of the learning rate and other hyperparameters
- <http://deeplearningbook.org>