


ConvNets For Self-Driving Cars



Code at: <https://github.com/kujason/avod>




Code at: <https://github.com/oandrienko/fast-semantic-segmentation>

Learning Objectives


- Learn how a neural network can use cross-correlation in its hidden layers instead of general matrix multiplication, to form ConvNets
- Learn the advantages of using ConvNets over traditional neural networks for processing images

ConvNets

- Used for processing data defined on grid
- 1D time series data, 2D images, 3D videos
- Two major type of layers:
 1. Convolution Layers
 2. Pooling Layers
- **Example: VGG 16**




Fully Connected VS Convolutional Layers



$$h_n = g(W^T h_{n-1} + b)$$

Fully Connected VS Convolutional Layers



$$h_n = g(W * h_{n-1} + b)$$

Cross Correlation

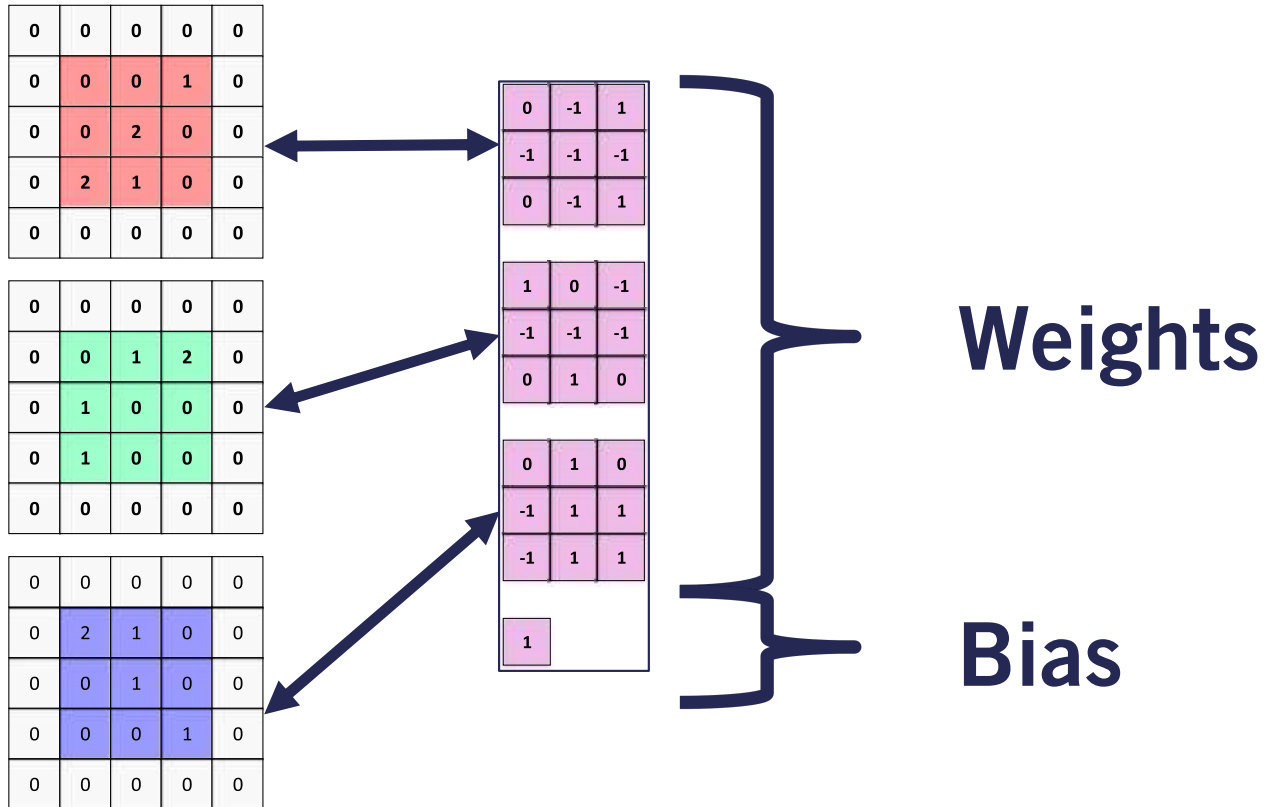
0	0	0	0	0
0	0	0	1	0
0	0	2	0	0
0	2	1	0	0
0	0	0	0	0

0	0	0	0	0
0	0	1	2	0
0	1	0	0	0
0	1	0	0	0
0	0	0	0	0


0	0	0	0	0
0	2	1	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0

- **Width:** horizontal dimension of input volume **3**
- **Height:** vertical dimension of input volume **3**
- **Depth:** number of channels of input volume **3**
- **Padding size:** essential to retain shape! **1**

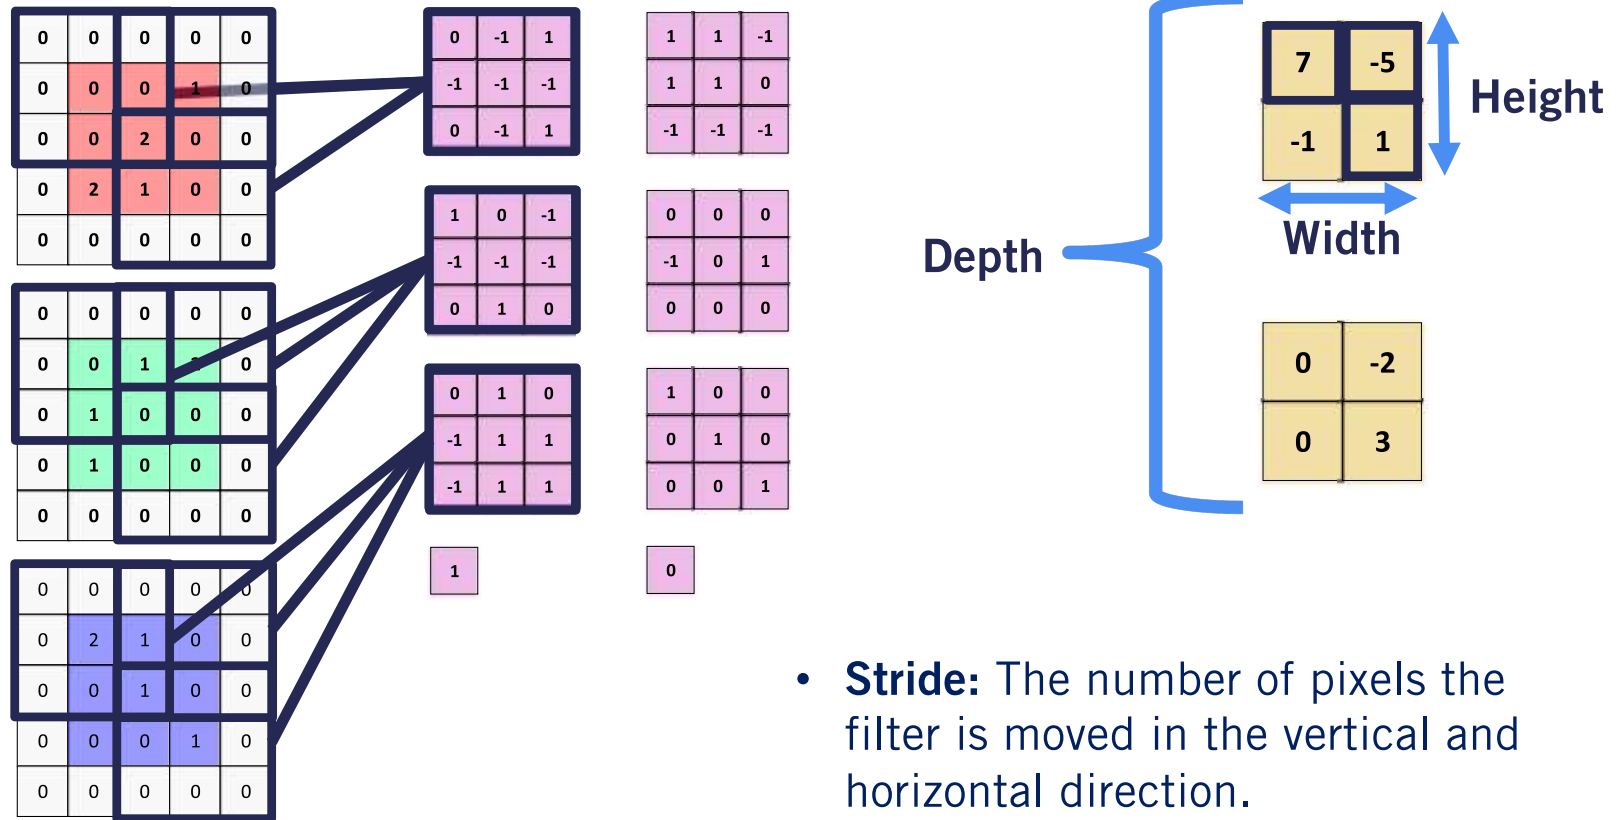
Cross Correlation



Cross Correlation



Cross Correlation



Output Volume Shape

- Filters are size $m \times m$
- Number of filters = K
- Stride = S , Padding = P

$$W_{out} = \frac{W_{in} - m + 2 \times P}{S} + 1$$

$$H_{out} = \frac{H_{in} - m + 2 \times P}{S} + 1$$

$$D_{out} = K$$

Pooling Layers: Max Pooling

21	8	8	12
12	19	9	7
8	10	4	3
18	9	10	9

$$\max(21, 8, 12, 19) = 21$$

21	12
18	10

Output Volume Shape

- Pool size $n \times n$
- Stride = S

$$W_{out} = \frac{W_{in} - n}{S} + 1$$

$$H_{out} = \frac{H_{in} - n}{S} + 1$$

$$D_{out} = D_{in}$$

Pooling Layers: Max Pooling

21	8	8	12
12	19	9	7
8	10	4	3
18	9	10	9

$$\max(21, 8, 12, 19) = 21$$

21	12
18	10

Advantages of ConvNets

- Convolutional neural networks are by design, a natural choice to process images
- Convolutional layers have **less parameters** than fully connected layers, reducing the chances of overfitting
- Convolutional layers use the same parameters to process every block of the image. Along with pooling layers, this leads to **translation invariance**, which is particularly important for image understanding

