

# Markov Decision Processes

# What is a Markov decision process?

- **MDP** (Environment evaluation)
- Andrey Markov (1856-1922)
- “Markov” generally means that given the **present state**, the **future** and **independent**
- $P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$


$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$



- A **mathematical** representation of a sequential decision making problem

# Markov Decision Process

- An MDP is defined by 5 tuple  $(S, A, \gamma, \{P_{sa}\}, R)$ 
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Probability that a from  $s$  leads to  $s'$
    - i.e.,  $P(s' | s, a)$
    - Also called the model
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state** (or distribution)
  - Maybe a **terminal state**
  - A **discount factor**:  $\gamma$
- MDPs are a family of non-deterministic search problems
  - Reinforcement learning: MDPs where we don't know the transition or reward functions



# Optimal Utilities

- Fundamental operation: compute the optimal values of states  $s$

- Optimal values ( $V^*(s)$ ) define optimal policies( $\pi^*(s)$ )

- Define the value of a state  $s$ :



$V^*(s)$  = expected utility starting in  $s$  and acting optimally

- Define the value of a q-state  $(s,a)$ :

$Q^*(s,a)$  = expected utility starting in  $s$ , taking action  $a$  and thereafter acting optimally

- Define the optimal policy:

$\pi^*(s)$  = optimal action from state  $s$



# Optimal Value function (Bellman equations)

- Given a fixed policy  $\pi$ , its value function  $V^\pi$  satisfies the

**Bellman equations:**

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P_{s\pi(s)}(s') V^\pi(s')$$

Immediate reward

expected sum of  
future discounted rewards

- We start in some state  $s_0$ , and get to choose some action  $a_0 \in A$
- As a result of our choice, the state of the MDP randomly transitions to some successor state  $s_1$ , drawn according to  $s_1 \sim P_{s_0 a_0}$
- Then, we get to pick another action  $a_1$

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$

## The Basic Setting for Learning

- **Training data:**  $n$  finite horizon trajectories, of the form  
 $\{s_0, a_0, r_0, \dots, s_T, a_T, r_T, s_{T+1}\}.$

- **Deterministic or stochastic policy:** A sequence of decision rules


$$\{\pi_0, \pi_1, \dots, \pi_T\}.$$

- Each  $\pi$  maps from the observable history (states and actions) to the action space at that time point. |

# Example of Q learning (episode 1)

- Initialize  $\hat{Q}$  to 0
- Random initial state =  $\langle 1, 1 \rangle$
- Random action from  $A_{\langle 1, 1 \rangle} = east$ 
  - $s' = \langle 2, 1 \rangle$
  - $R_a(s, s') = 0$
- Update  $\hat{Q}(\langle 1, 1 \rangle, east) = 0$
- Random action from  $A_{\langle 2, 1 \rangle} = north$ 
  - $s' = \langle 2, 0 \rangle$
  - $R_a(s, s') = 100$
- Update  $\hat{Q}(\langle 2, 1 \rangle, north) = 100$
- No more moves possible, start again...

$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$



# Example of Q learning

$$\hat{Q}(s, a) = R_a(s, s') + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')$$
$$\gamma = 0.5$$

- Random Initial State  $\langle 0,0 \rangle$
- Update  $\hat{Q}(\langle 1,1 \rangle, east) = 50$
- Update  $\hat{Q}(\langle 1,2 \rangle, east) = 25$

